

# BRAIN STROKE DETECTION

Using Support Vector  
Machines



## 1. Cleaning data

In [244]:

```
import pandas as pd
from google.colab import drive

drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

In [245]:

```
raw_data=pd.read_csv('/content/drive/MyDrive/datasets/stroke_data.csv')
raw_data.head()
```

Out[245]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67.0	0	1	Yes	Private	Urban
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural
2	31112	Male	80.0	0	1	Yes	Private	Rural
3	60182	Female	49.0	0	0	Yes	Private	Urban
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural

In [246]:

```
raw_data.describe()
```

Out[246]:

	id	age	hypertension	heart_disease	avg_glucose_level	brain_atrial_fibrillation
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893200
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854000
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000

In [247]:

```
raw_data.isnull().sum()
```

Out[247]:

```
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              201
smoking_status    0
stroke            0
dtype: int64
```

In [248]:

```
import numpy as np

from sklearn.impute import SimpleImputer

imputer=SimpleImputer()

raw_data['bmi']=imputer.fit_transform(np.array(raw_data['bmi']).reshape(-1, 1))
```

In [249]:

```
raw_data.isnull().sum()
```

Out[249]:

```
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              0
smoking_status    0
stroke            0
dtype: int64
```

In [250]:

```
clean_data=raw_data.drop('id',axis=1)
```

In [251]:

```
clean_data.head()
```

Out[251]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg.
0	Male	67.0	0	1	Yes	Private	Urban	
1	Female	61.0	0	0	Yes	Self-employed	Rural	
2	Male	80.0	0	1	Yes	Private	Rural	
3	Female	49.0	0	0	Yes	Private	Urban	
4	Female	79.0	1	0	Yes	Self-employed	Rural	



In [252]:

```

import matplotlib.pyplot as plt

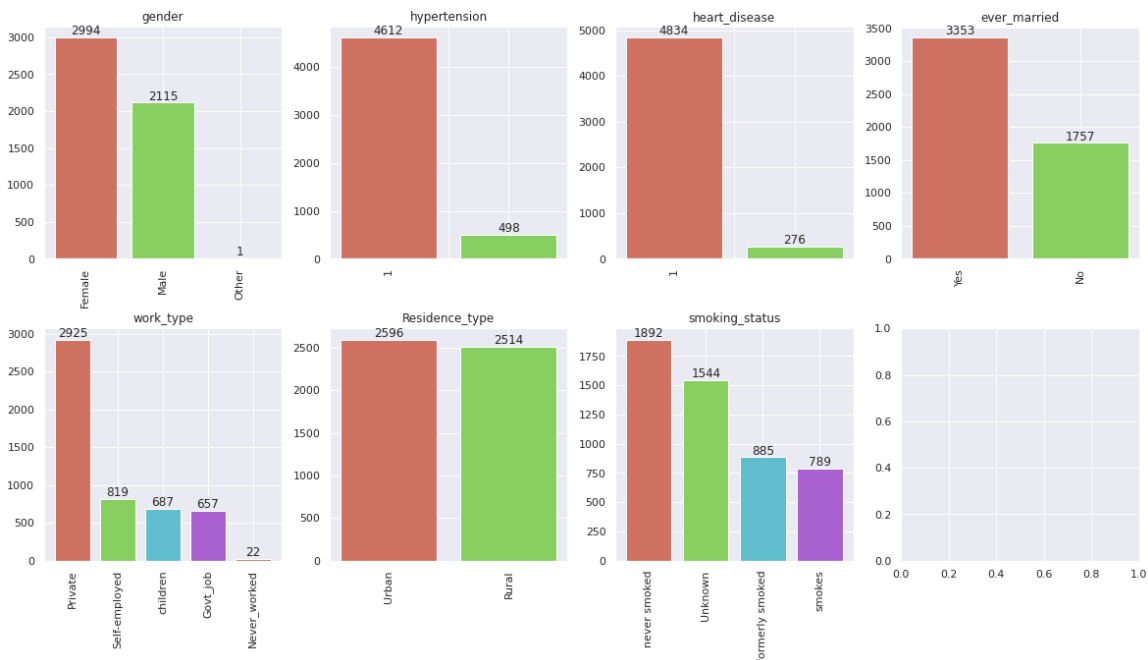
features=clean_data.columns.tolist()
features.remove('stroke')

con_features=['age','bmi','avg_glucose_level']
cat_features=[feature for feature in features if feature not in con_features]

fig, axes = plt.subplots(2, 4)
plt.subplots_adjust(hspace=0.3)
fig.set_figwidth(20)
fig.set_figheight(10)
feature_index=0
colors=['#CF7261','#87CF61','#61BECF','#A961CF']
for i,ai in enumerate(axes):
    for j,aij in enumerate(axes[i]):
        if(not (i == 1 and j == 3)):
            value_counts=clean_data[cat_features[feature_index]].value_counts()
            value_counts_index=value_counts.index.tolist()
            value_counts_values=value_counts.values.tolist()
            bars=axes[i][j].bar(value_counts_index,value_counts_values,color=colors)
            axes[i][j].set_title(cat_features[feature_index])
            axes[i][j].set_xticklabels(value_counts_index,rotation=90)
            feature_index+=1

            for k,bar in enumerate(bars):
                axes[i][j].text(k,value_counts_values[k], value_counts_values[k],ha='center',va='bottom')

```



In work type, both children and never worked convey same meaning with regard to employment status. So both categories are merged into one as never worked

In [253]:

```
clean_data['work_type']=clean_data['work_type'].replace('children', 'Never_worked')
```

## 2. Encoding categorical variables

In [254]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

label=['ever_married', 'Residence_type']
one_hot=['gender', 'work_type', 'smoking_status']

for feature in cat_features:
    if(feature in label):
        encoder=LabelEncoder()
        clean_data[feature]=encoder.fit_transform(np.array(clean_data[feature]).reshape(-1,1))
    elif(feature in one_hot):
        dummies=pd.get_dummies(clean_data[feature], prefix=feature)
        clean_data=pd.concat([clean_data, dummies], axis=1)
        clean_data=clean_data.drop([feature], axis=1)

new_features=clean_data.columns.tolist()
new_features.remove('stroke')

cat_features=[feature for feature in new_features if feature not in con_features]

clean_data.head()
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:11
5: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
  y = column_or_1d(y, warn=True)
```

Out[254]:

	age	hypertension	heart_disease	ever_married	Residence_type	avg_glucose_level	
0	67.0	0	1	1	1	228.69	36.60
1	61.0	0	0	1	0	202.21	28.89
2	80.0	0	1	1	0	105.92	32.50
3	49.0	0	0	1	1	171.23	34.40
4	79.0	1	0	1	0	174.12	24.00

## 3. Splitting testing data

In [255]:

```

from sklearn.model_selection import train_test_split as tts

x=clean_data.drop('stroke',axis=1)
y=clean_data['stroke']

x_train,x_test,y_train,y_test=tts(x,y,test_size=0.2,random_state=1)

print(x_train.shape,x_test.shape,y_train.shape,y_test.shape)
print(y_train.value_counts())
print(y_test.value_counts())

```

```

(4088, 18) (1022, 18) (4088,) (1022,)
0      3899
1       189
Name: stroke, dtype: int64
0       962
1        60
Name: stroke, dtype: int64

```

## 4. Feature selection

In [256]:

```

features=clean_data.drop('stroke',axis=1)
target=clean_data['stroke']

```

## ANOVA test

In [257]:

```

from sklearn.feature_selection import f_classif

anova_result=pd.DataFrame(columns=['feature','f_score','p_value'])

for feature in con_features:
    result = f_classif(np.array(x_train[feature]).reshape(-1,1),y_train)
    anova_result=anova_result.append({'feature':feature,'f_score':round(result[0][0],5),
    },'p_value':round(result[1][0],5)},ignore_index=True)

anova_result.sort_values(by=['f_score'],ascending=False,ignore_index=True)

```

Out[257]:

	feature	f_score	p_value
0	age	240.88241	0.0000
1	avg_glucose_level	57.13297	0.0000
2	bmi	6.13381	0.0133

## Chi - Squared test

In [258]:

```

from sklearn.feature_selection import chi2

chi2_result=pd.DataFrame(columns=['feature','chi2_score','p_value'])

for feature in cat_features:
    result = chi2(np.array(x_train[feature]).reshape(-1,1),y_train)
    chi2_result=chi2_result.append({'feature':feature,'chi2_score':round(result[0][0],5),
    'p_value':round(result[1][0],5)},ignore_index=True)

chi2_result.sort_values(by=['chi2_score'],ascending=False,ignore_index=True)

```

Out[258]:

	feature	chi2_score	p_value
0	heart_disease	67.66219	0.00000
1	hypertension	57.51814	0.00000
2	work_type_Never_worked	24.17382	0.00000
3	ever_married	12.10966	0.00050
4	work_type_Self-employed	10.57817	0.00114
5	smoking_status_formerly smoked	10.41427	0.00125
6	smoking_status_Unknown	7.13372	0.00756
7	Residence_type	0.68596	0.40754
8	work_type_Private	0.54293	0.46122
9	gender_Male	0.21359	0.64397
10	gender_Female	0.15344	0.69527
11	smoking_status_never smoked	0.04851	0.82568
12	work_type_Govt_job	0.00033	0.98546
13	smoking_status_smokes	0.00004	0.99481
14	gender_Other	NaN	NaN

In [259]:

```

x_train=x_train.drop(['work_type_Govt_job','Residence_type','gender_Male','gender_Other',
'gender_Female','smoking_status_never smoked'],axis=1)

```

## 5. Handling imbalanced data



In [260]:

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()

print(y_train.value_counts())
sns.countplot(y_train)
```

0 3899

1 189


Name: stroke, dtype: int64

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[260]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f63298eed50>



In [261]:

```
from imblearn.over_sampling import RandomOverSampler

sampler = RandomOverSampler(random_state=1)

x_resampled, y_resampled = sampler.fit_resample(x_train, y_train)
```

In [262]:

```
print(y_resampled.value_counts())  
  
sns.countplot(y_resampled)
```

0 3899

1 3899


Name: stroke, dtype: int64

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[262]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f6324e28a50>



## 6. Modeling

In [270]:

```
from sklearn.svm import SVC  
from sklearn.metrics import accuracy_score  
  
model=SVC(C=100,gamma='auto')  
model.fit(x_resampled,y_resampled)  
  
y_train_pred=model.predict(x_resampled)  
  
print('Training accuracy {}'.format(round(accuracy_score(y_resampled,y_train_pred)*100,2)))
```

Training accuracy 99.99%

Removing weak features from test set

In [264]:

```
x_test=x_test.drop(['work_type_Govt_job','Residence_type','gender_Male','gender_Other',  
'gender_Female','smoking_status_never smoked'],axis=1)
```

Make predictions for test set

In [265]:

```
y_test_pred=model.predict(x_test)  
pd.Series(y_test_pred).value_counts()
```

Out[265]:

```
0    992  
1     30  
dtype: int64
```

In [271]:

```
print('Test accuracy {}'.format(round(accuracy_score(y_test,y_test_pred)*100,2)))
```

Test accuracy 92.37%

In [267]:

```
# !jupyter nbconvert --to html "/content/drive/MyDrive/Colab Notebooks/svm.ipynb"
```